

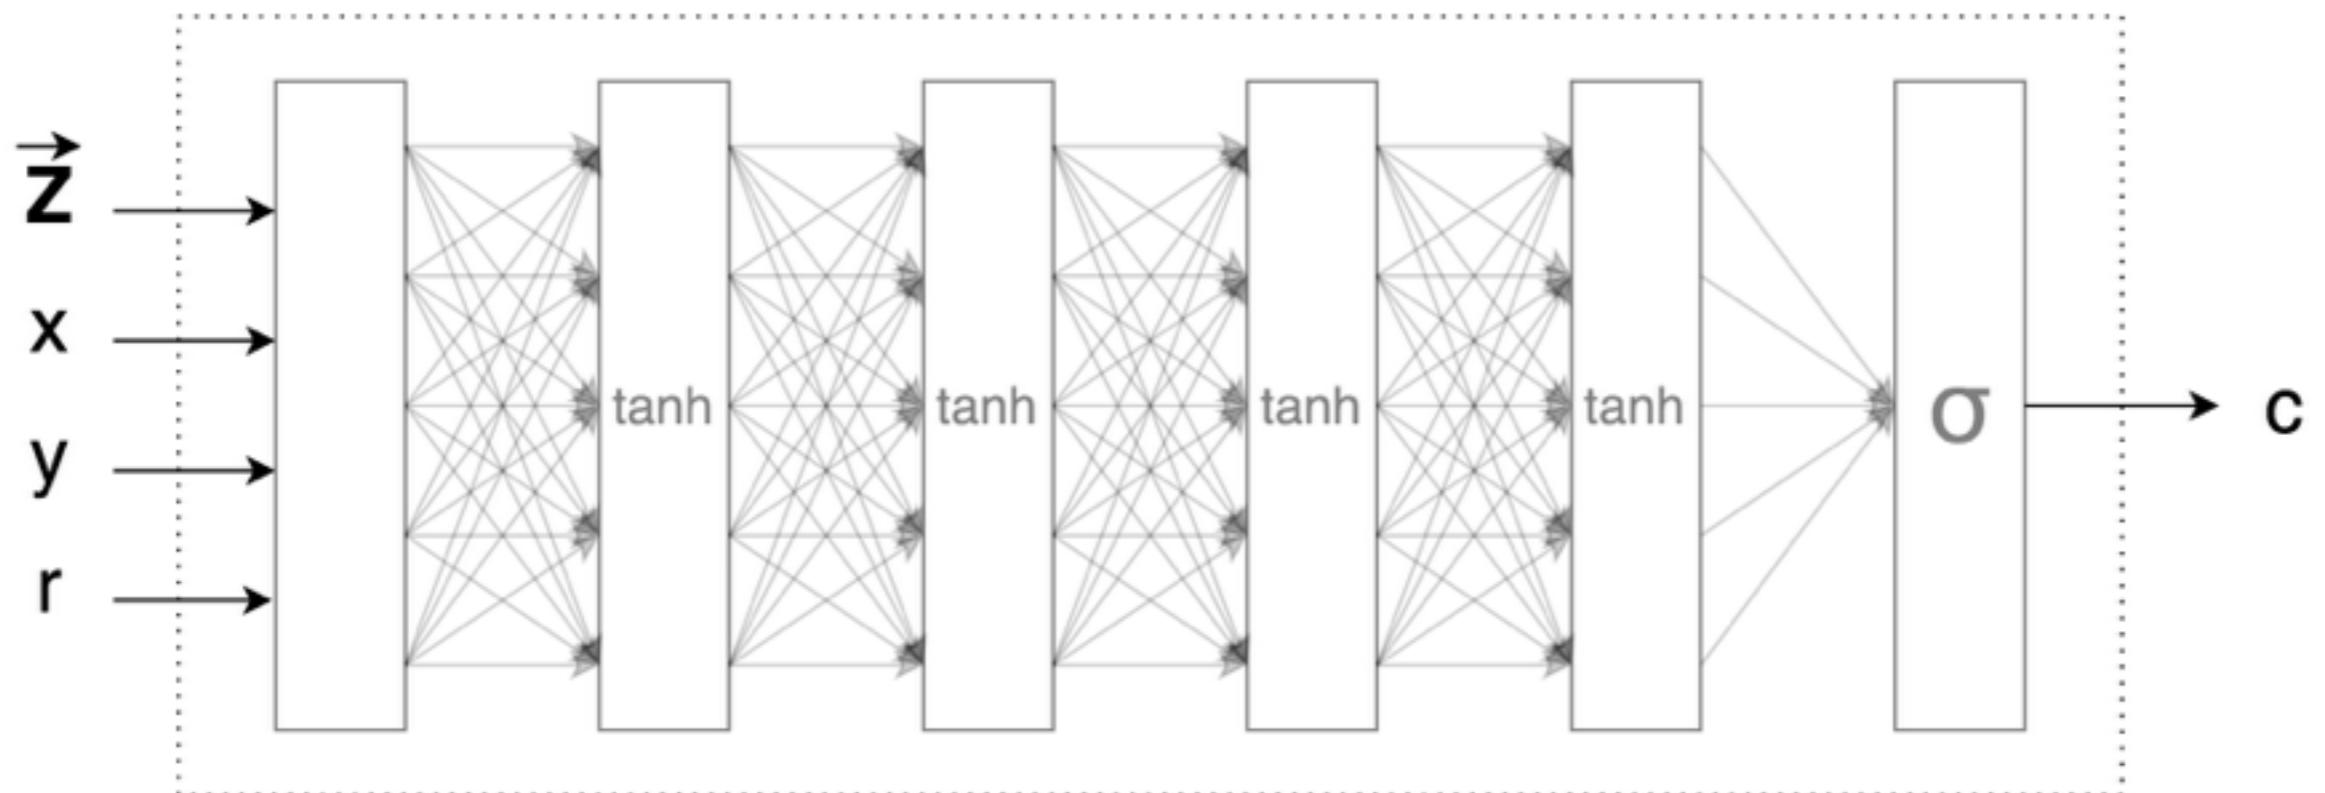
9. 更多的网络类型

丰富的神经网络类型

- 深度学习除了经典的卷积网络、循环网络之外还有广泛的网络类型，应用在不同的问题。

一个简单但是有趣的例子

- CPPN (Compositional Pattern Producing Network)



- 网络输入是像素的坐标值 (x,y)
- 网络输出是像素的RGB值
- 把 $[0,0] - [100,100]$ 坐标逐个输入， 将输出的RGB值组成完整图像， 将会是什么样子？



怎么实现的？

```
1 require 'torch'
2 require 'image'
3 require 'nn'

4

5

6 local ConvLayer = function(in_size, out_size)
7     local conv = nn.SpatialConvolutionMM(in_size, out_size, 1, 1)
8     conv.weight:normal(0, 1.0)

9

10    return conv
11 end

12

13 local buildNet = function(input_size, output_size, net_number, layer_number)
14     local net = nn.Sequential()

15

16     net:add( ConvLayer(input_size, net_number) )
17     net:add( nn.Tanh() )
18     for i = 1, layer_number - 1 do
19         net:add( ConvLayer(net_number, net_number) )
20         net:add( nn.Tanh() )
21     end
22     net:add( ConvLayer(net_number, output_size) )
23     net:add( nn.Sigmoid() )

24

25     return net
26 end
```

```
local buildXYRInput = function(width, height)
    local xyr_in = torch.Tensor(3, height, width)

    for i = 1, width do
        xyr_in[{1, 1, i}] = (i - width/2)
    end
    for i = 2, height do
        xyr_in[{1, i, {}}]:copy( xyr_in[{1, 1, {}}] )
    end

    for i = 1, height do
        xyr_in[{2, i, 1}] = (i - height/2)
    end
    for i = 2, width do
        xyr_in[{2, {}, i}]:copy( xyr_in[{2, {}, 1}] )
    end

    xyr_in[3] = torch.cmul(xyr_in[1], xyr_in[1] ) + torch.cmul(xyr_in[2], xyr_in[2])
    xyr_in[3]:sqrt()

    return xyr_in
end

local net = buildNet(3, 3, 32, 5)
local xyr = buildXYRInput(640, 640)
xyr = xyr / 500.0
local out = net:forward(xyr)
image.savePNG('./output.png',out)
```



<http://blog.otoro.net/2016/03/25/generating-abstract-patterns-with-tensorflow/>

孪生网络(Siamese)

- 孪生网络(Siamese network)是一种网络结构，通过一个NN将样本的维度降低到某个较低的维度。
- 在低维空间，任意两个样本：
 - 如果它们是相同类别，空间距离尽量接近0
 - 如果它们是不同类别，空间距离大于某个间隔

https://github.com/Teaonly/easyLearning.io/tree/master/siamese_network

```
+++++
HingeEmbeddingCriterion
```

```
    criterion = nn.HingeEmbeddingCriterion([margin])
```

Creates a criterion that measures the loss given an input x which is a 1-dimensional vector and a label y (1 or -1).

This is usually used for measuring whether two inputs are similar or dissimilar, e.g. using the L1 pairwise distance, and is typically used for learning nonlinear embeddings or semi-supervised learning.

```
        |  $x_i$ , if  $y_i == 1$ 
loss( $x, y$ ) =  $1/n \{$ 
        |  $\max(0, \text{margin} - x_i)$ , if  $y_i == -1$ 
```

If x and y are n -dimensional Tensor s, the sum operation still operates over all the elements, and divides by n (this can be avoided if one sets the internal variable sizeAverage to false). The margin has a default value of 1 , or can be set in the constructor.

```
+++++
```

```
1 require 'nn'
2
3 function build_model()
4
5     local cnn = nn.Sequential()
6     cnn:add(nn.SpatialConvolution(1, 8, 3, 3, 1, 1))
7     cnn:add(nn.ReLU())
8     cnn:add(nn.SpatialMaxPooling(2, 2))
9     cnn:add(nn.SpatialConvolution(8, 16, 3, 3, 1, 1))
10    cnn:add(nn.ReLU())
11    cnn:add(nn.SpatialMaxPooling(2, 2))
12    cnn:add(nn.SpatialConvolution(16, 16, 3, 3, 1, 1))
13    cnn:add(nn.ReLU())
14    cnn:add(nn.SpatialMaxPooling(2, 2))
15    cnn:add(nn.SpatialConvolution(16, 32, 3, 3, 1, 1))
16    cnn:add(nn.ReLU())
17
18    cnn:add(nn.Reshape(32*4*4))
19    cnn:add(nn.Linear(32*4*4, 512))
20    cnn:add(nn.ReLU())
21    cnn:add(nn.Linear(512, 2))
22
23    --The siamese model
24    --clone the encoder and share the weight, bias. Must also share the gradWeight and gradBias
25    siamese_encoder = nn.ParallelTable()
26    siamese_encoder:add(cnn)
27    siamese_encoder:add(cnn:clone('weight','bias', 'gradWeight','gradBias'))
28
29    --The siamese model (inputs will be Tensors of shape (2, channel, height, width))
30    model = nn.Sequential()
31    model:add(nn.SplitTable(2)) -- split input tensor along the rows (1st dimension) to table for input to ParallelTable
32    model:add(siamese_encoder)
33    model:add(nn.PairwiseDistance(2)) --L2 pariwise distance
34
35    return model
36 end
```

```

13 local loadBatch = function(index, isTest)
14   local ii = index or 1
15
16   local data = g.allTrainSamples.data
17   local labels = g.allTrainSamples.labels
18   if ( isTest ) then
19     data = g.allTestSamples.data
20     labels = g.allTestSamples.labels
21   end
22
23   local number = data:size(1)
24
25   local batch = {}
26   batch.data = torch.Tensor(g.batchSize, 2, 1, 32, 32)
27   batch.label = torch.zeros(g.batchSize) - 1
28   for i = 1, g.batchSize do
29     batch.data[i][1]:copy ( data[ii] )
30     local iii = g.siamese[ii]
31     batch.data[i][2]:copy ( data[iii] )
32
33     if ( labels[ii] == labels[iii] ) then
34       batch.label[i] = 1
35     end
36
37     ii = ii + 1
38     if ( ii > number ) then
39       ii = 1
40     end
41   end
42
43   if _CUDA_ then
44     batch.data = batch.data:cuda()
45     batch.label = batch.label:cuda()
46   end
47
48   return batch, ii
49 end

```

```

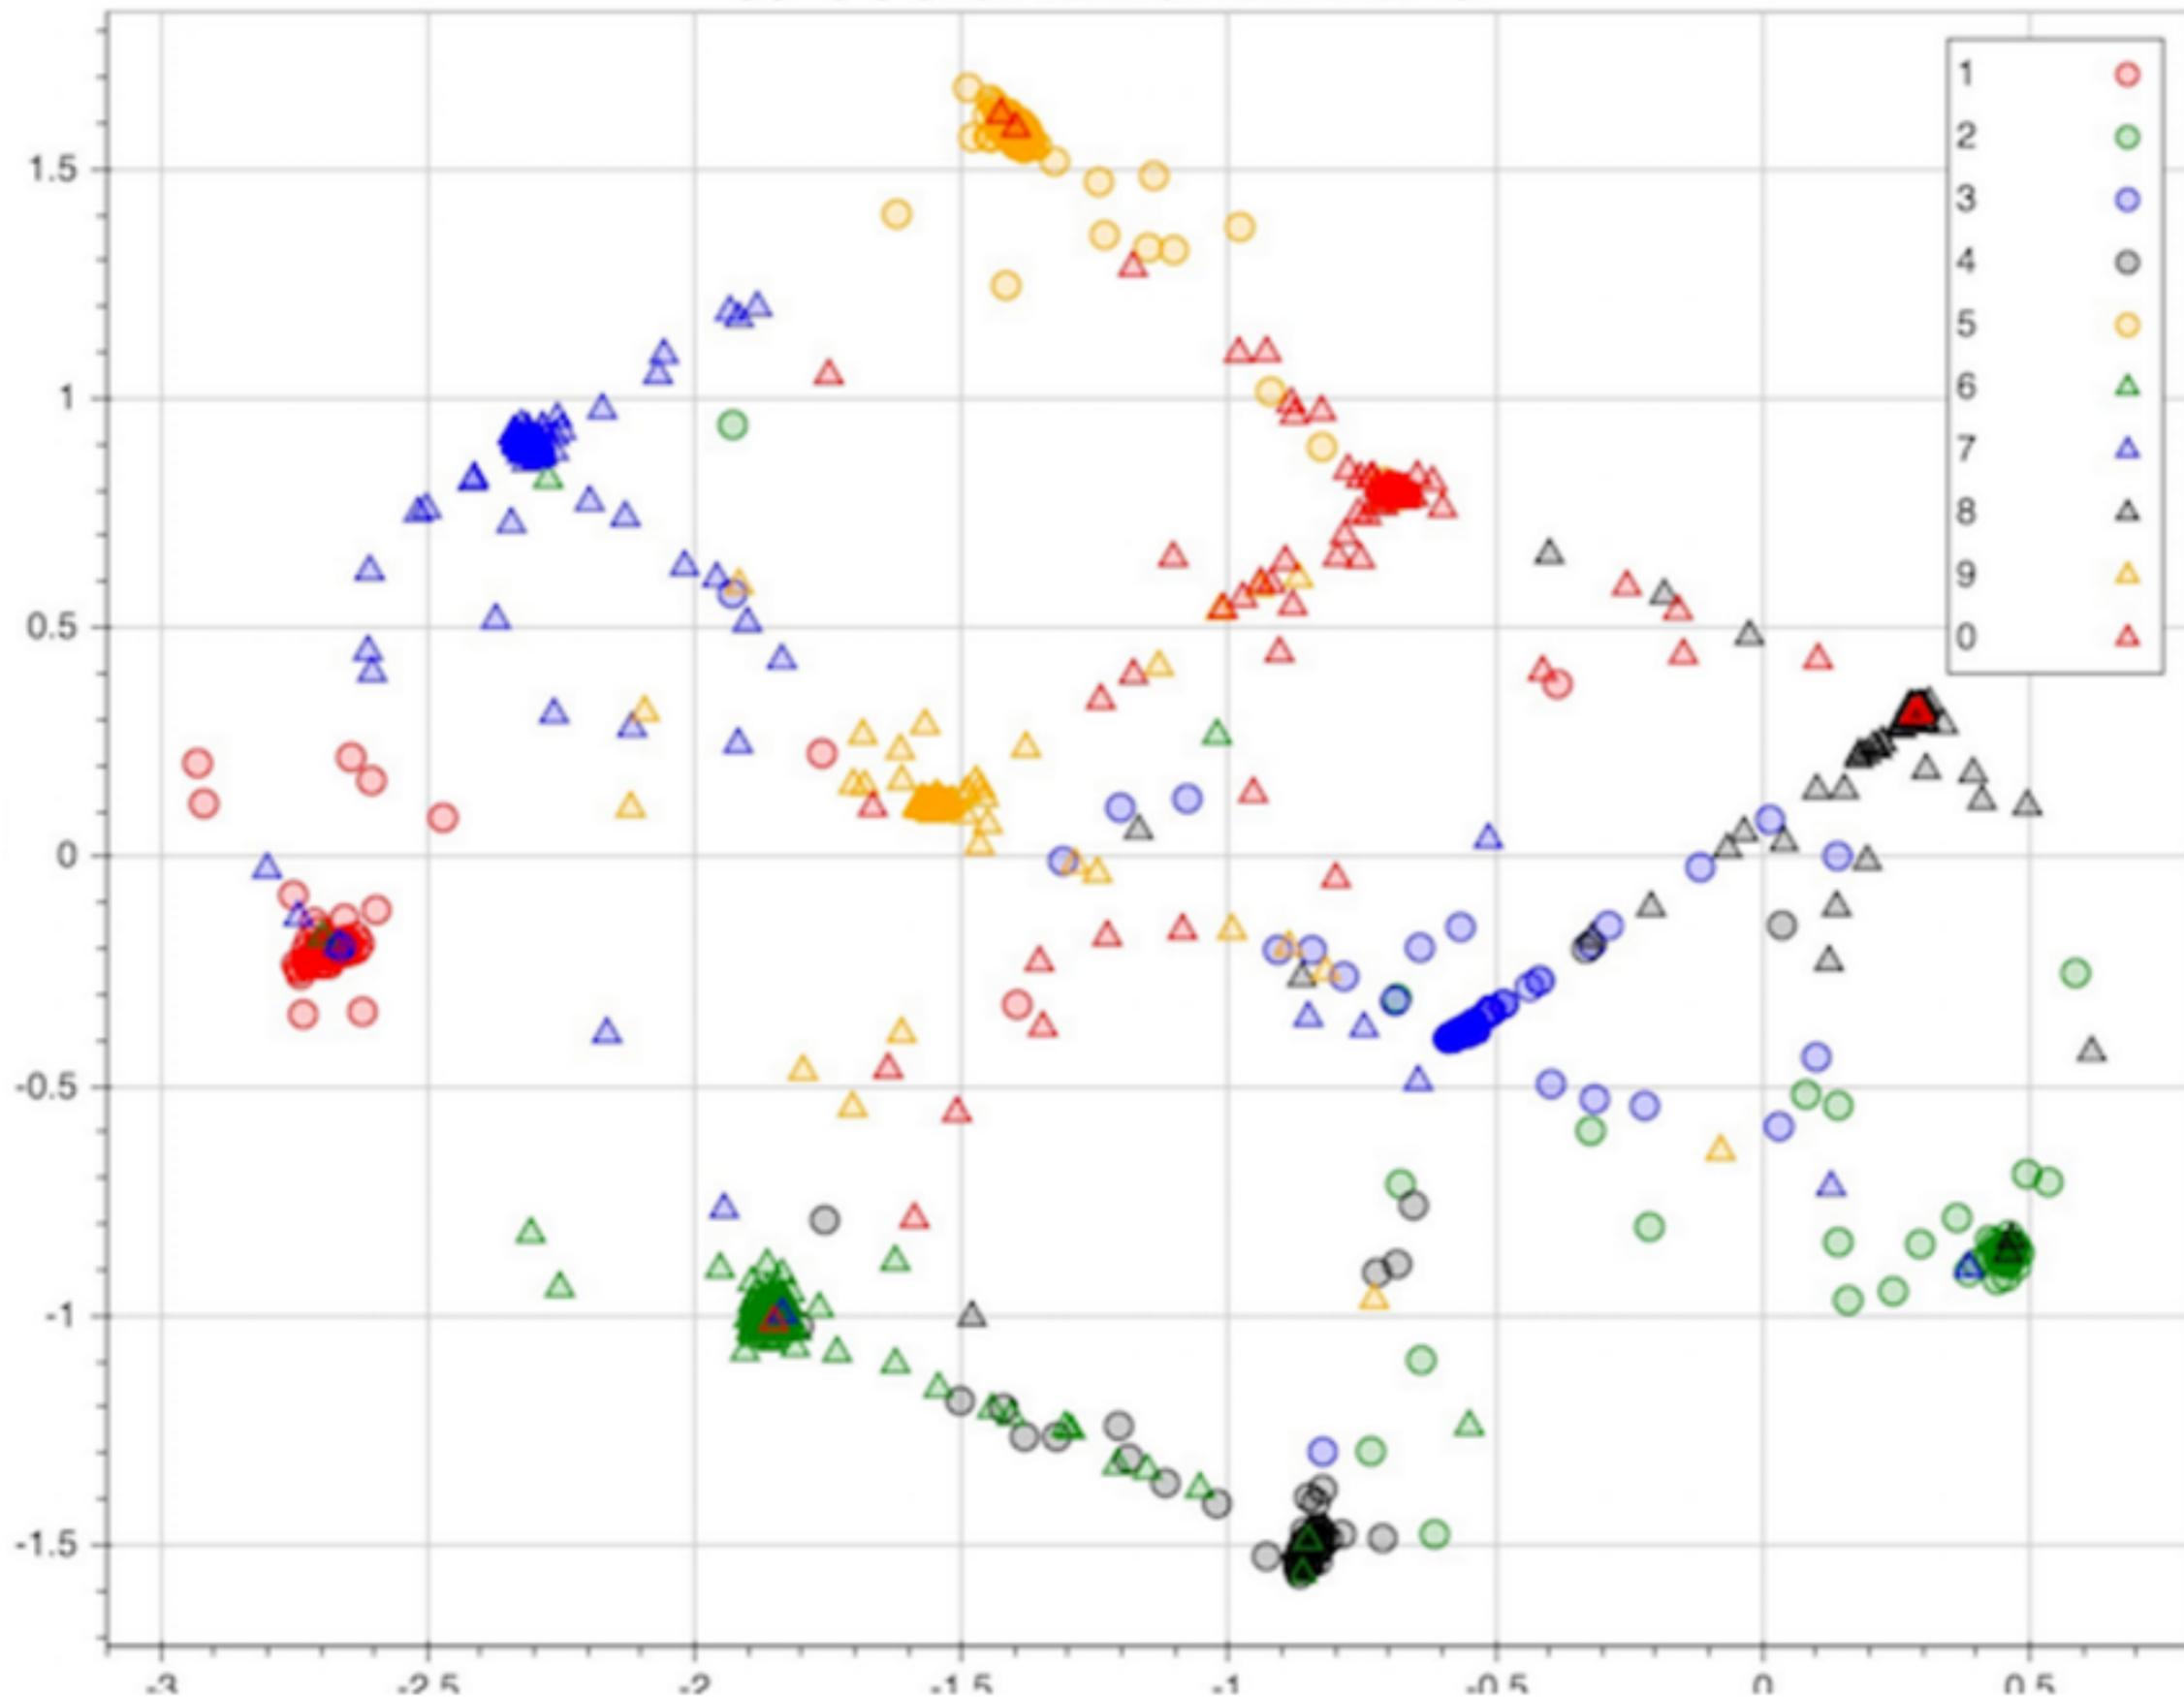
51   local doTrain = function()
52     print(">>>>>>>>>TRAINING>>>>>>>>>");  

53
54     if _CUDA_ then
55       g.model:cuda()
56       g.criterion:cuda()
57     end
58     g.model:training()  

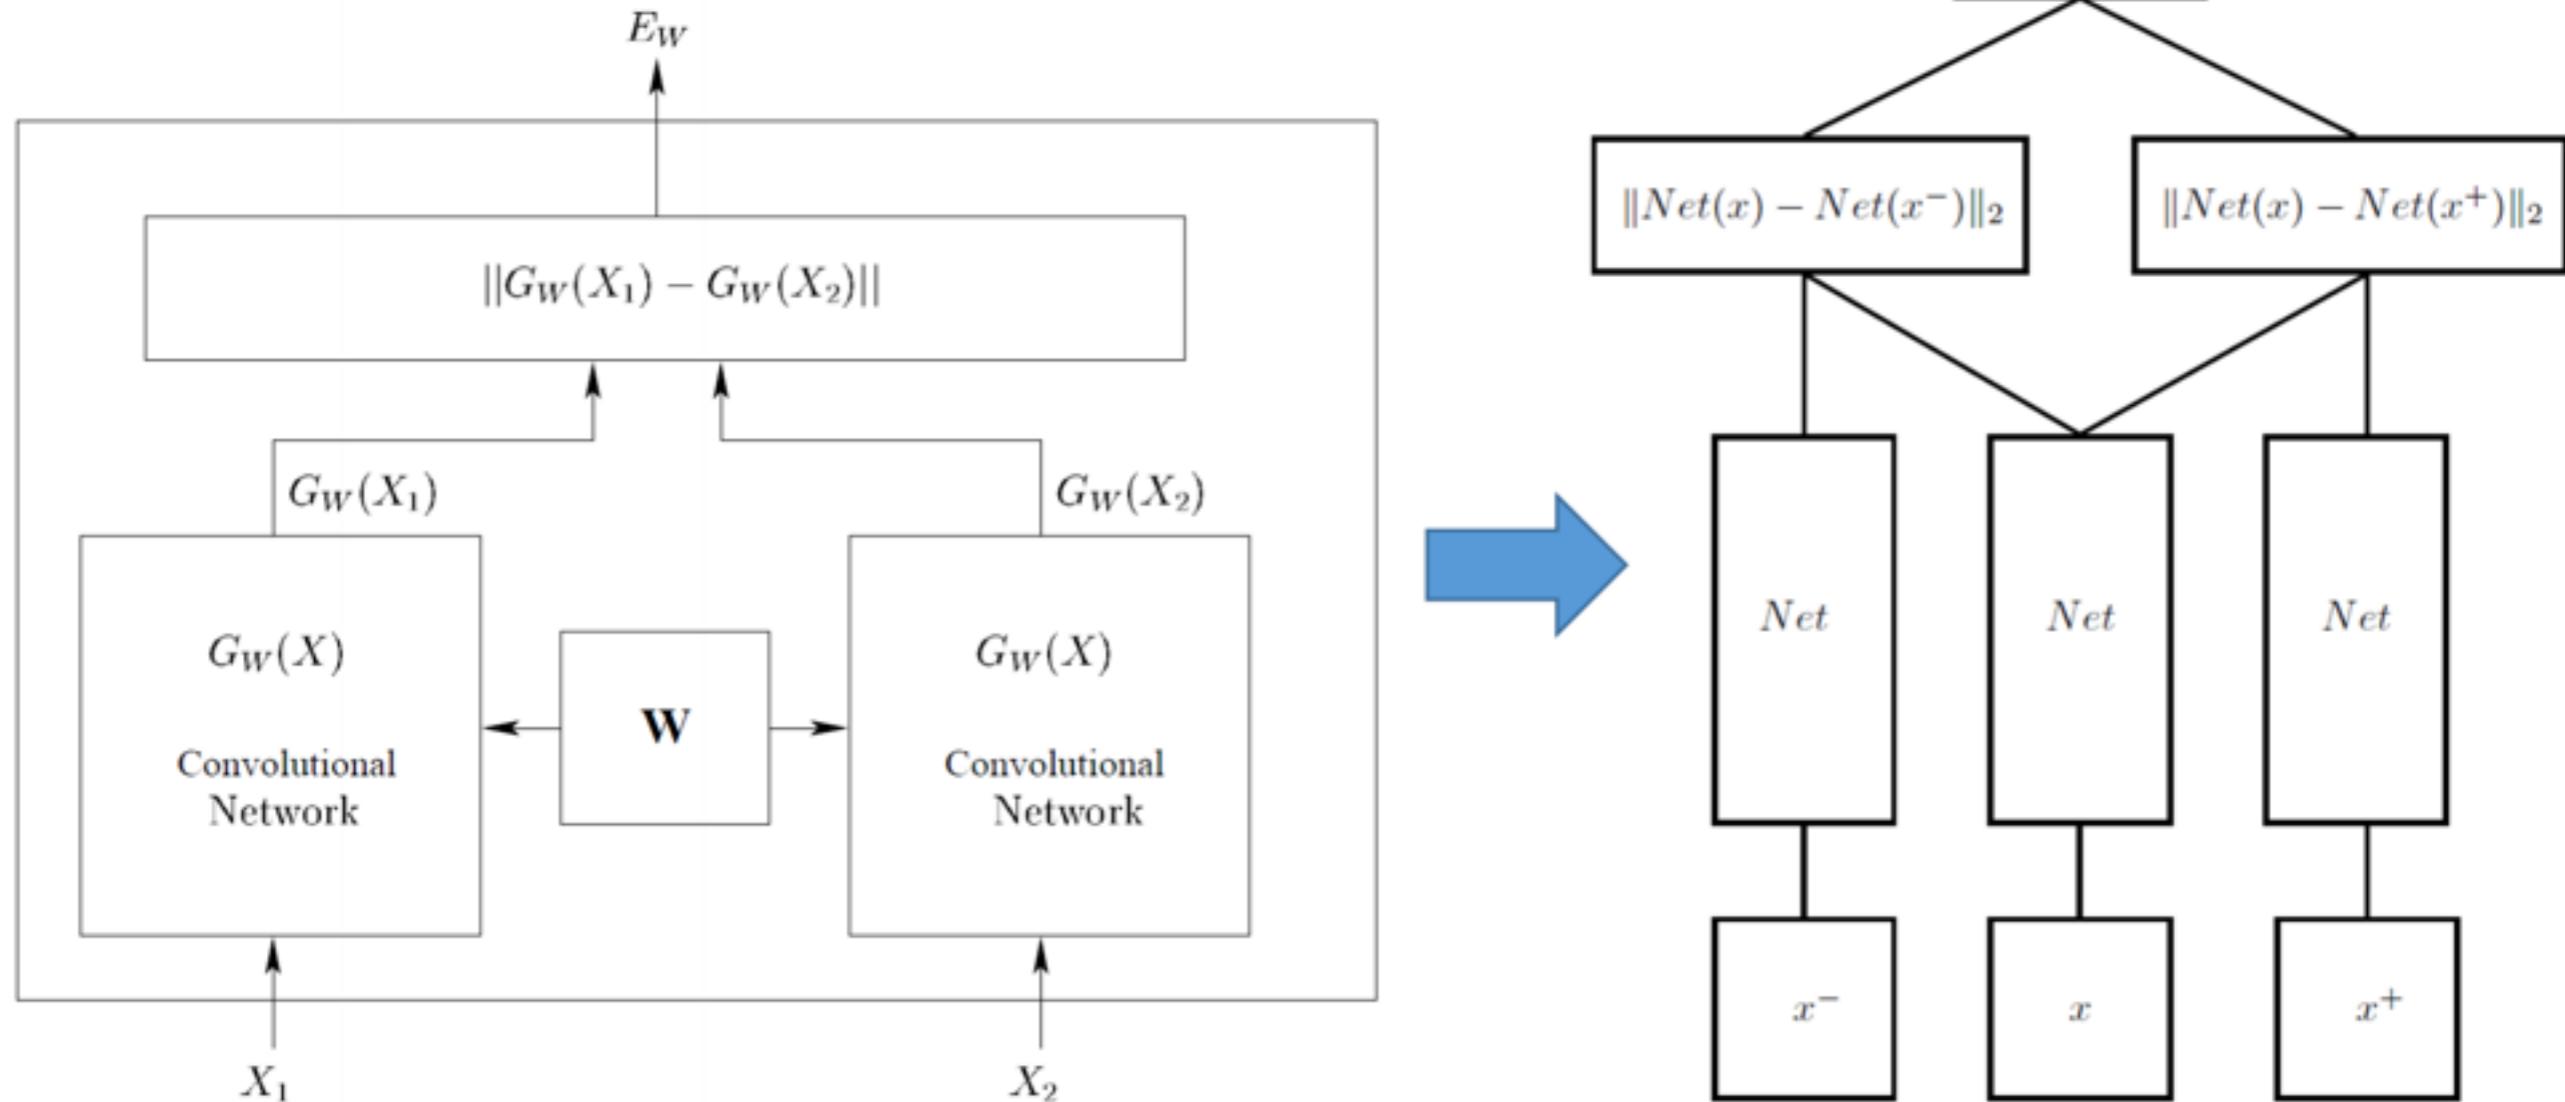
59
60     local parameters,gradParameters = model:getParameters()
61     local batch = nil
62     local feval = function(x)
63       -- get new parameters
64       if x ~= parameters then
65         parameters:copy(x)
66       end
67
68       -- reset gradients
69       gradParameters:zero()
70
71       local distance = g.model:forward(batch.data)
72
73       local f = g.criterion:forward(distance, batch.label)
74       local df = g.criterion:backward(distance, batch.label)
75
76       g.model:backward(batch.data, df)
77
78       return f, gradParameters
79     end
80
81     local index = 1
82     local maxIterate = torch.floor( g.allTrainSamples.data:size(1) / g.batchSize )
83     for i = 1, maxIterate do
84       batch, index = loadBatch(index)
85
86       g.optim(feval, parameters, g.optimState)
87
88       collectgarbage();
89       xlua.progress(i, maxIterate)
90     end
91   end

```

样本降维到2维时的分布



Triplet Network



From Siamese to Triplet Network



- 应用类型

- Image ranking
- Face verification
- Metric learning

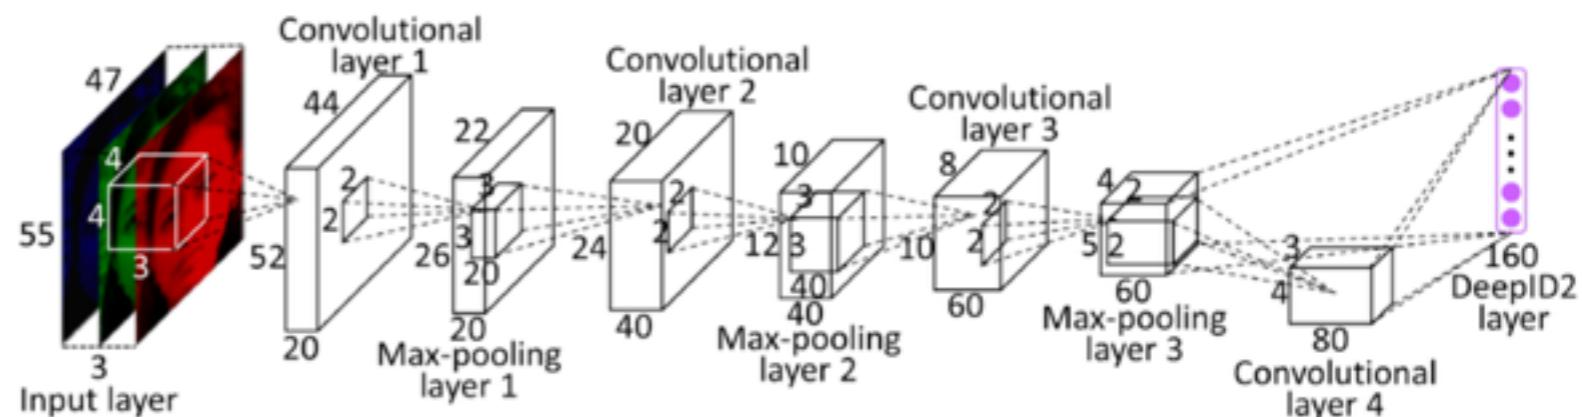
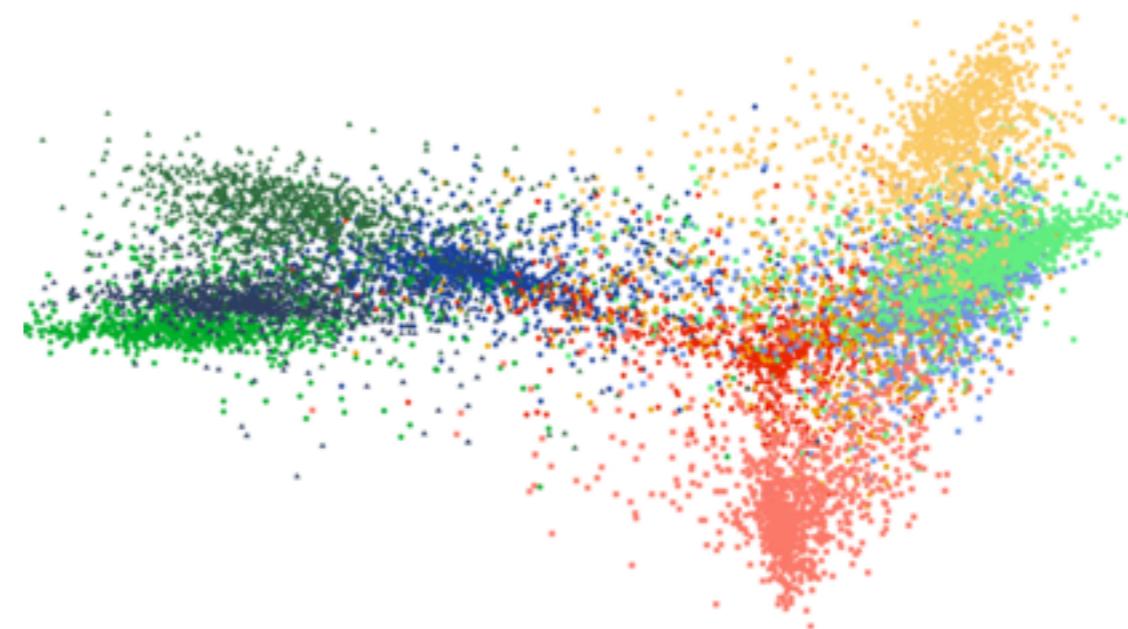


Figure 1: The ConvNet structure for DeepID2 extraction.



Variational Auto-encoder

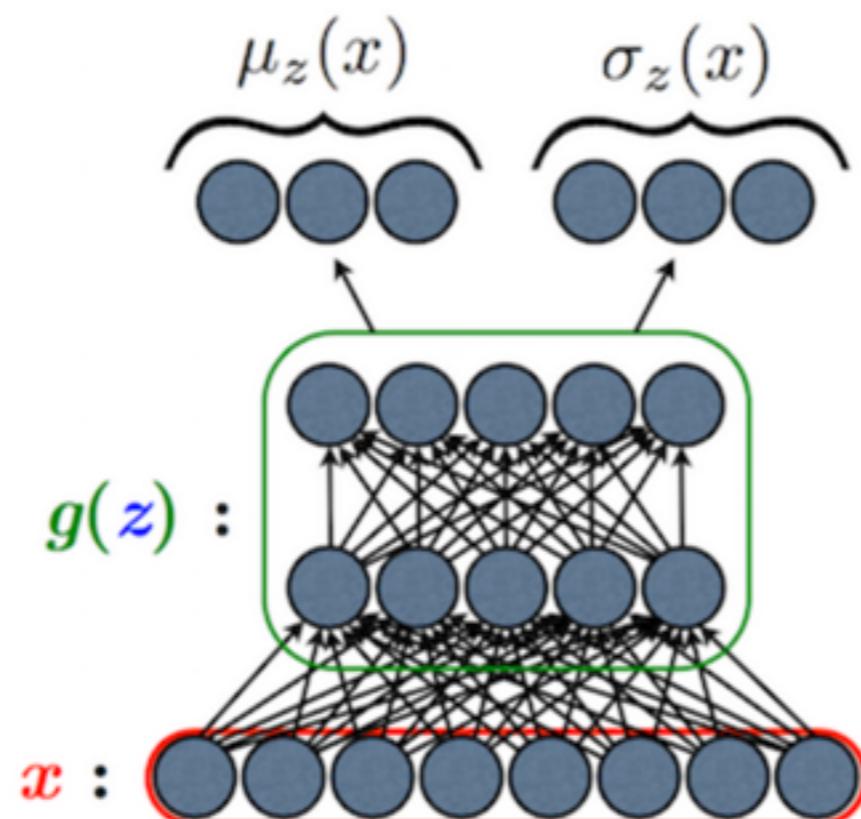
- 将Variational Bayes引入到神经网络中
- 网络和经典的Auto Encoder类似
 - 编码网络：将输入编码到一个参数（隐变量）分布
 - 解码网络：对分布的参数进行采样，拟合输入数据

首先引入隐变量 z , 利用“编码网络”拟合参数化的后验概率 $q_\phi(z|x)$

这里的 ϕ 就是神经网络的权重, 网络输出为隐变量 z 的条件分布参数 (即后验概率), 这里采用的是高斯分布:

$$q_\phi(z|x) = N(z; \mu_z(x), \sigma_z(x))$$

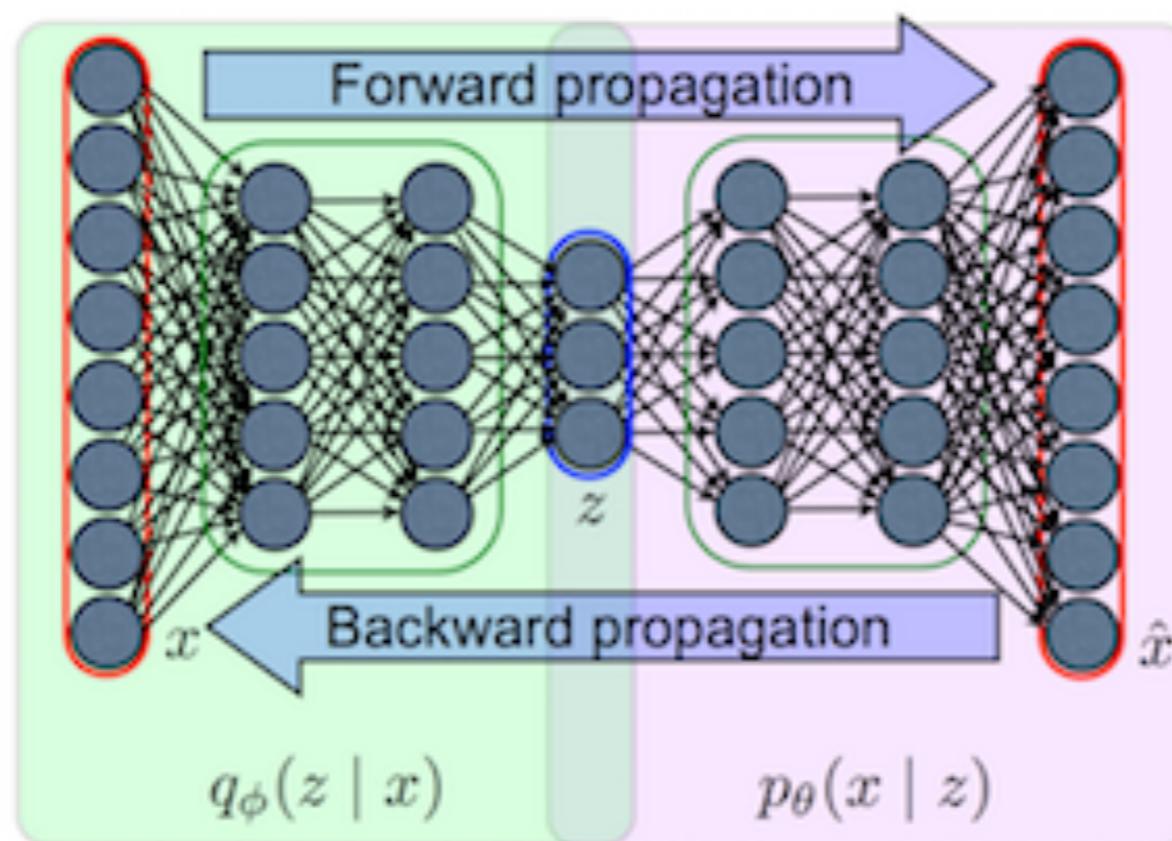
其中 $\mu_z(x), \sigma_z(x)$, 就是Encoder网络的输出



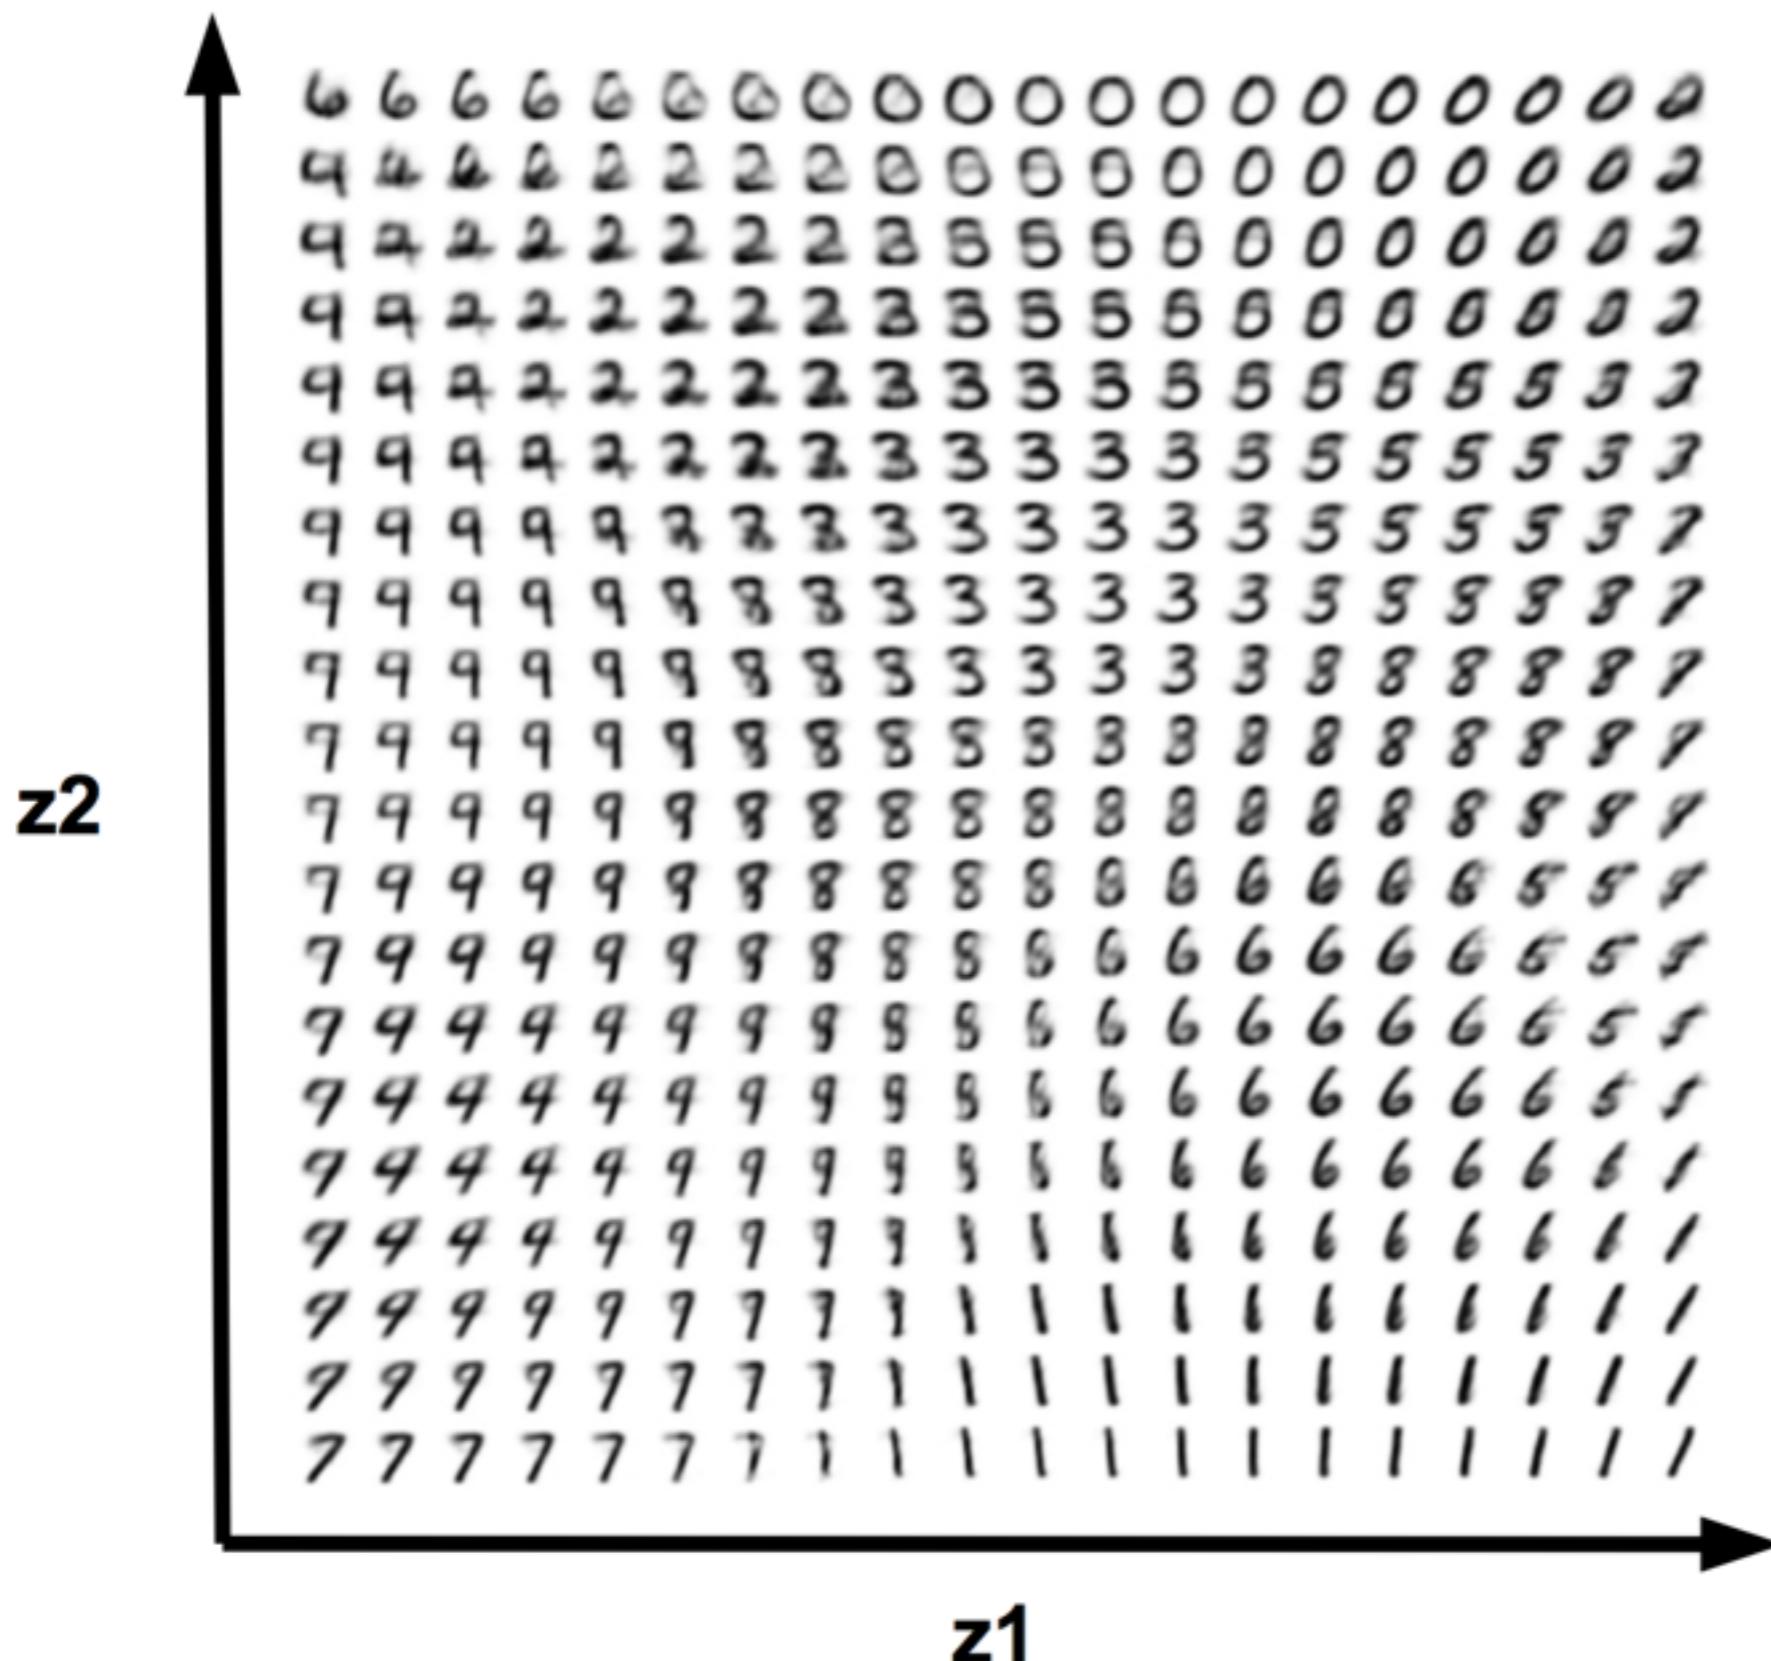
2. 解码网络和变分推断

解码网络是已知隐变量 z , 计算样本 x 的条件概率, 这个概率是似然概率, 参数为 θ , 即计算: $p_\theta(x|z)$

$$\text{Objective function: } \mathcal{L}(\theta, \phi, x) = -D_{\text{KL}}(q_\phi(z|x) \| p_\theta(z)) + \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)]$$



2D Latent space: MNIST



强化学习中的Policy network

- Markov decision processes
- Bellman equations
- Policy network & Policy Gradient Methods

Markov decision processes

- S : 系统所处状态的集合
- A : 系统可以做出的动作的集合
- P_{sa} 状态转移的概率，即在状态 s 下采取动作 a 之后，新状态的概率
- $\gamma \in [0, 1)$ 折现系数
- 奖励函数 $R : S \times A \mapsto \mathbb{R}$ $R : S \mapsto \mathbb{R}$

- 状态序列

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

- 某个过程的价值函数定义 (重视当前的奖励, 遥远未来的奖励会逐步降低)

$$R(s_0, a_0) + \gamma R(s_1, a_1) + \gamma^2 R(s_2, a_2) + \dots \quad \text{与 } a \text{ 有关}$$

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots \quad \text{与 } a \text{ 无关}$$

- 系统目标, 找到一组动作, 过程奖励期望最大

$$\mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

- 决定动作输出的策略 (policy) 定义为一个函数：

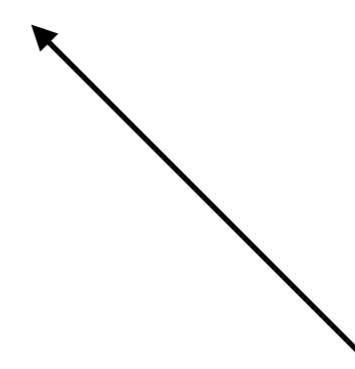
$$\pi : S \mapsto A$$

- 在某个策略下，过程的价值函数

$$V^\pi(s) = \mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi].$$

- Bellman 公式：

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s').$$



 $\pi : S \mapsto A$

- 最优价值函数可以定义为：

$$V^*(s) = \max_{\pi} V^\pi(s).$$

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s').$$

- 最优的决策函数，可以定义为：

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

注意最优决策函数， 和初始状态s无关

$$V^*(s) = V^{\pi^*}(s) \geq V^\pi(s).$$

- 价值迭代，求解 $V^*(s)$

1. For each state s , initialize $V(s) := 0$.
2. Repeat until convergence {

For every state, update $V(s) := R(s) + \max_{a \in A} \gamma \sum_{s'} P_{sa}(s')V(s')$

}

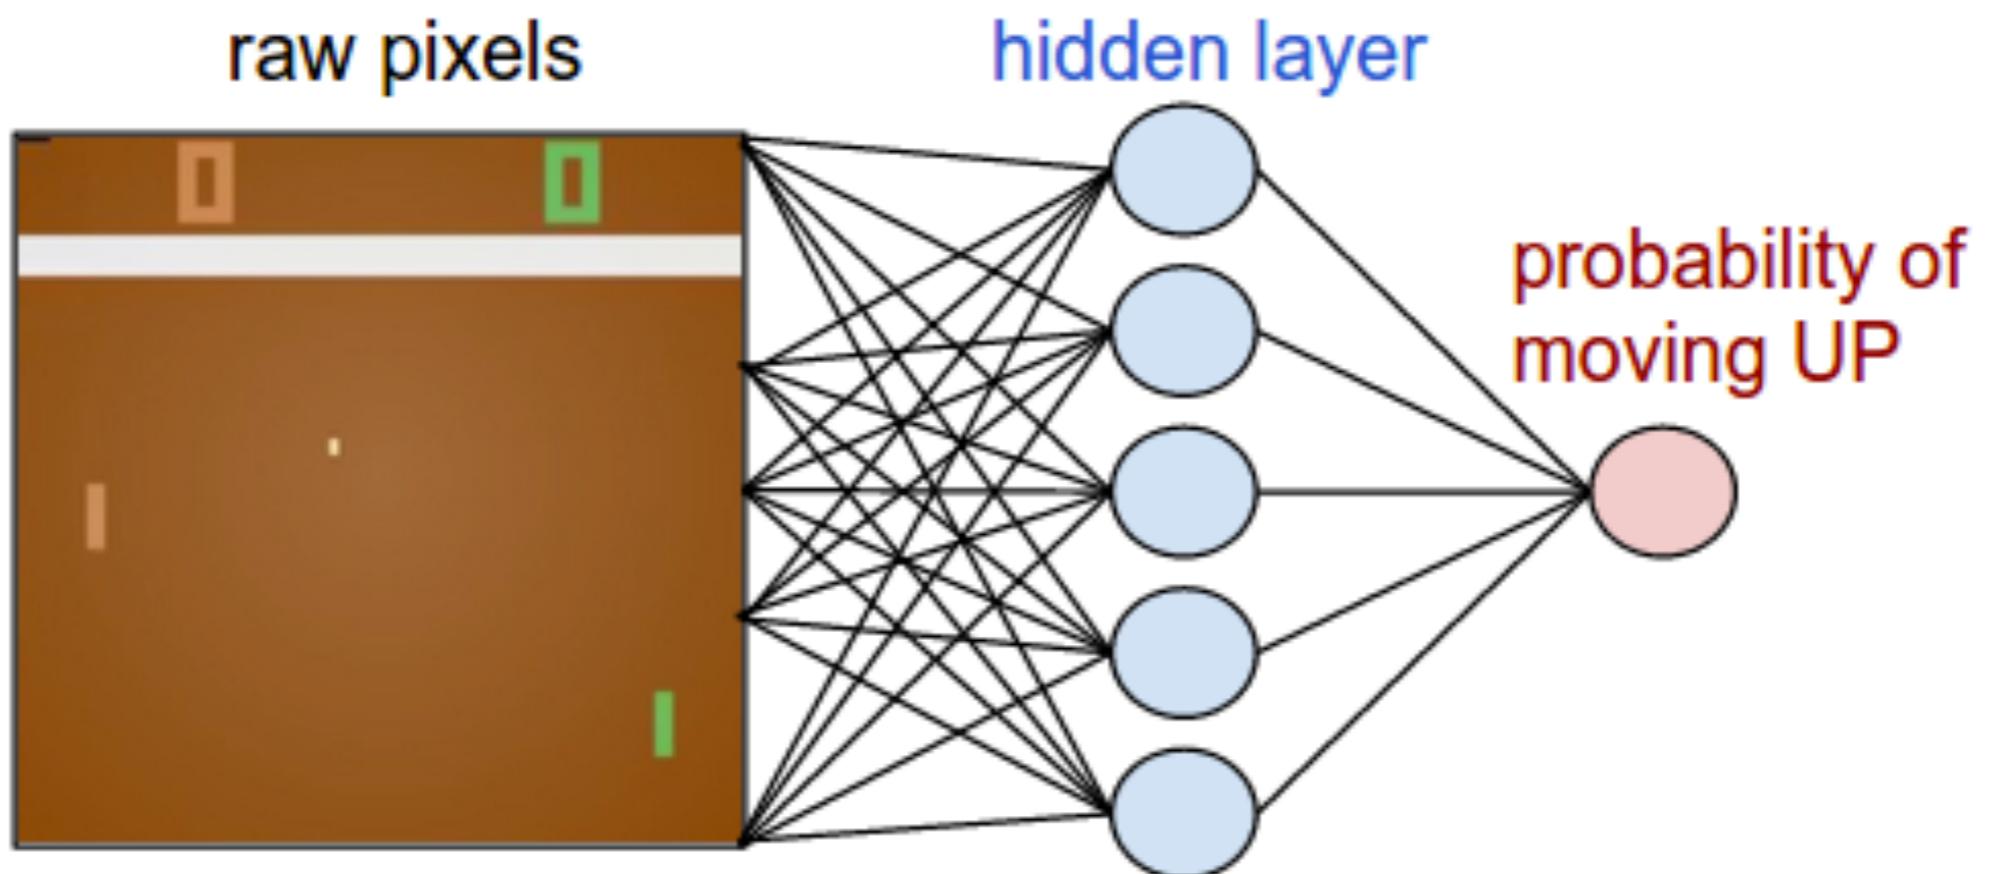
- 决策迭代，直接求解 $\pi^*(s)$

1. Initialize π randomly.
2. Repeat until convergence {
 - (a) Let $V := V^\pi$.
 - (b) For each state s , let $\pi(s) := \arg \max_{a \in A} \sum_{s'} P_{sa}(s')V(s')$

}

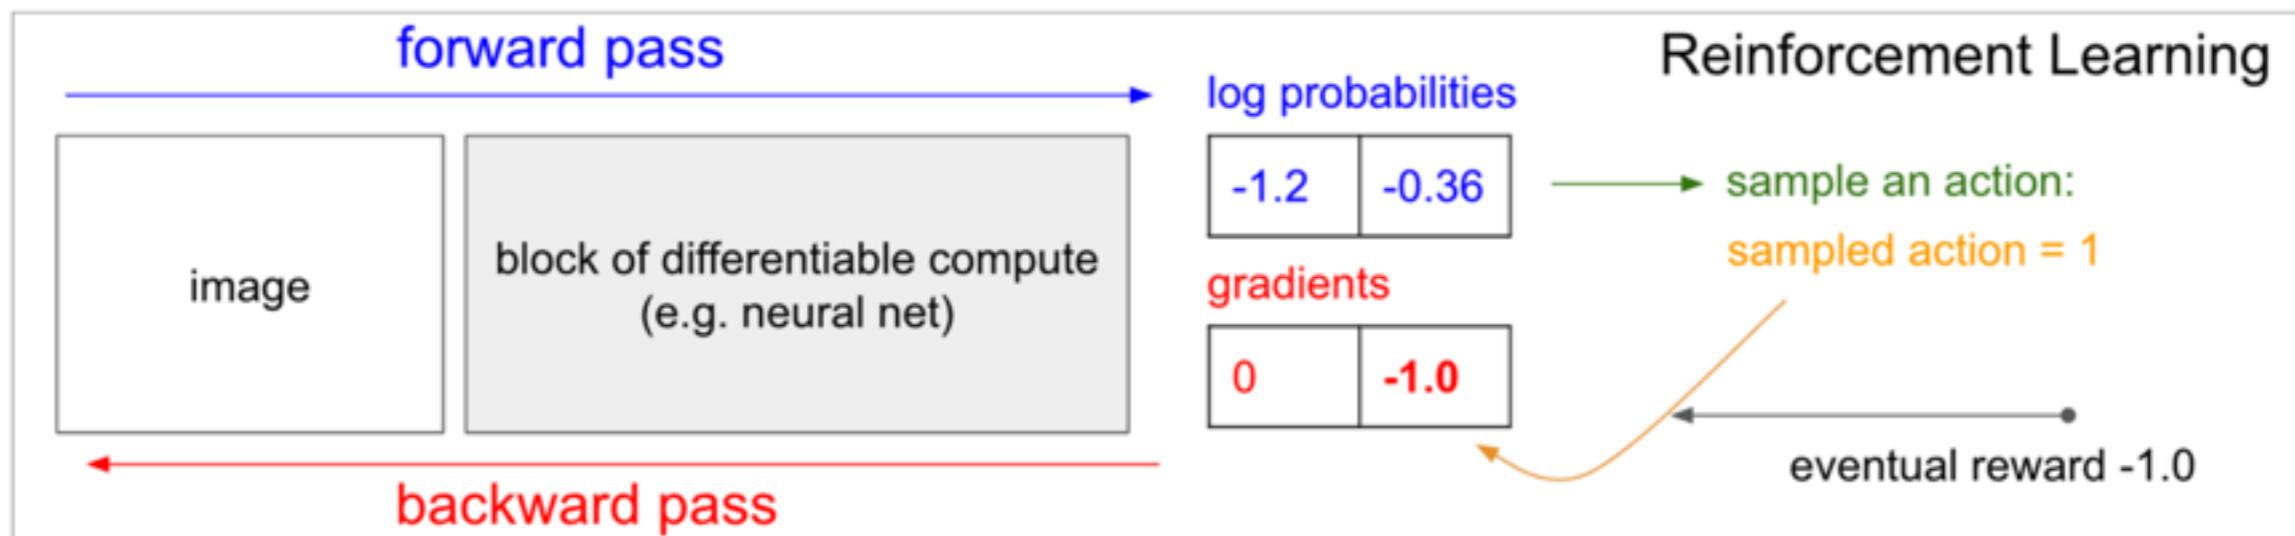
$|S| < \infty, |A| < \infty$

Policy network

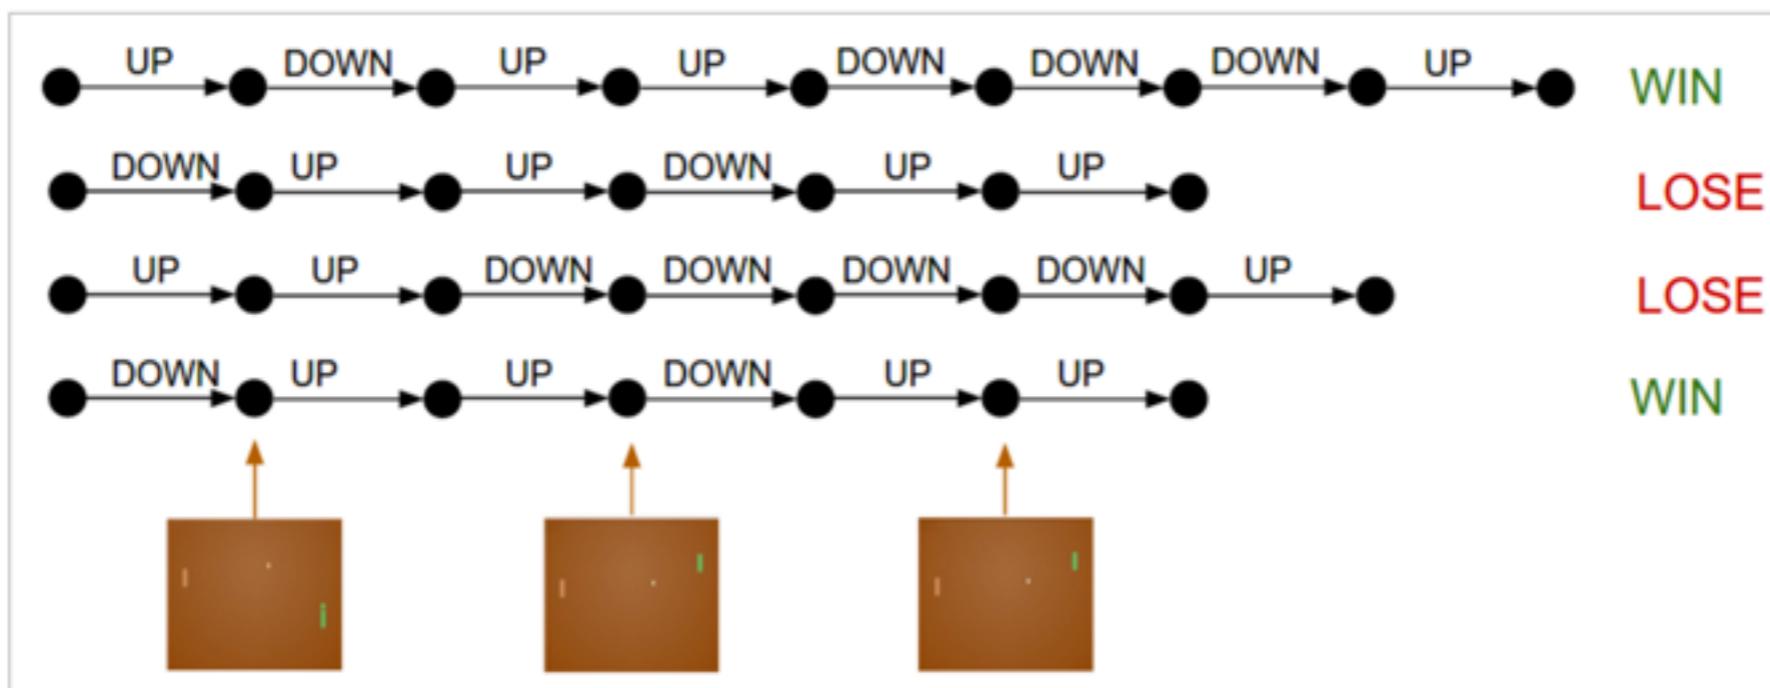


- 通过采样Policy Network, 取得当前状态的输出

- Policy Gradient Methods



Policy Gradients: Run a policy for a while. See what actions led to high rewards. Increase their probability.



- 1. 随机初始化Policy network
- 2. 进行大量迭代 {
 - 固定network，执行多次网络，得到一系列结果。
 - 通过BP算法，增大好的结果概率，减少差的结果概率
 - }

原理：通过多次采样平均梯度方向，逐步得需要的期望

- 实现一个简单的Policy Network

- 使用OpenAI Gym工具



```
0: game finished, reward: -1.000000
0: game finished, reward: -1.000000
resetting env. episode reward total was -21.000000. running mean: -21.000000
1: game finished, reward: -1.000000
1: game finished, reward: -1.000000 !!!!!!!
1: game finished, reward: -1.000000
resetting env. episode reward total was -20.000000. running mean: -20.990000
2: game finished, reward: -1.000000
```

<https://gym.openai.com/docs>

<http://karpathy.github.io/2016/05/31/r/>

1.准备模型，定义参数

```
1  """ Trains an agent with (stochastic) Policy Gradients on Pong. Uses OpenAI Gym. """
2  import numpy as np
3  import cPickle as pickle
4  import gym
5
6  # hyperparameters
7  H = 200 # number of hidden layer neurons
8  batch_size = 10 # every how many episodes to do a param update?
9  learning_rate = 1e-4
10 gamma = 0.99 # discount factor for reward
11 decay_rate = 0.99 # decay factor for RMSProp leaky sum of grad^2
12 resume = False # resume from previous checkpoint?
13 render = False
14
15 # model initialization
16 D = 80 * 80 # input dimensionality: 80x80 grid
17 if resume:
18     model = pickle.load(open('save.p', 'rb'))
19 else:
20     model = {}
21     model['W1'] = np.random.randn(H,D) / np.sqrt(D) # "Xavier" initialization
22     model['W2'] = np.random.randn(H) / np.sqrt(H)
23
24 grad_buffer = { k : np.zeros_like(v) for k,v in model.iteritems() } # update buffers that add up gradients over a batch
25 rmsprop_cache = { k : np.zeros_like(v) for k,v in model.iteritems() } # rmsprop memory
```

2.辅助函数

```
27 def sigmoid(x):
28     return 1.0 / (1.0 + np.exp(-x)) # sigmoid "squashing" function to interval [0,1]
29
30 def prepro(I):
31     """ prepro 210x160x3 uint8 frame into 6400 (80x80) 1D float vector """
32     I = I[35:195] # crop
33     I = I[::2,::2,0] # downsample by factor of 2
34     I[I == 144] = 0 # erase background (background type 1)
35     I[I == 109] = 0 # erase background (background type 2)
36     I[I != 0] = 1 # everything else (paddles, ball) just set to 1
37     return I.astype(np.float).ravel()
38
39 def discount_rewards(r):
40     """ take 1D float array of rewards and compute discounted reward """
41     discounted_r = np.zeros_like(r)
42     running_add = 0
43     for t in reversed(xrange(0, r.size)):
44         if r[t] != 0: running_add = 0 # reset the sum, since this was a game boundary (pong specific!)
45         running_add = running_add * gamma + r[t]
46         discounted_r[t] = running_add
47     return discounted_r
```

3. Forward/Backward函数

```
49 def policy_forward(x):
50     h = np.dot(model['W1'], x)
51     h[h<0] = 0 # ReLU nonlinearity
52     logp = np.dot(model['W2'], h)
53     p = sigmoid(logp)
54     return p, h # return probability of taking action 2, and hidden state
55
56 def policy_backward(eph, epdlogp):
57     """ backward pass. (eph is array of intermediate hidden states) """
58     dW2 = np.dot(eph.T, epdlogp).ravel()
59     dh = np.outer(epdlogp, model['W2'])
60     dh[eph <= 0] = 0 # backpro prelu
61     dW1 = np.dot(dh.T, epx)
62     return {'W1':dW1, 'W2':dW2}
--
```

4. 训练 (1)

```
64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprob, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprob else 3 # roll the dice!
```

4.训练 (2)

```
83 # record various intermediates (needed later for backprop)
84 xs.append(x) # observation
85 hs.append(h) # hidden state
86 y = 1 if action == 2 else 0 # a "fake label"
87 dlogps.append(y - aprob) # grad that encourages the action that was taken to be taken (see http://cs231n.gi
88
89 # step the environment and get new measurements
90 observation, reward, done, info = env.step(action)
91 reward_sum += reward
92
93 drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
--
```

注意我们保存了每一步的x, h,以及reward, 为bp计算准备

4.训练 (3)

```
98     # stack together all inputs, hidden states, action gradients, and rewards for this episode
99     epx = np.vstack(xs)
100    eph = np.vstack(hs)
101    epdlogp = np.vstack(dlogps)
102    epr = np.vstack(drs)
103    xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105    # compute the discounted reward backwards through time
106    discounted_epr = discount_rewards(epr)
107    # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108    discounted_epr -= np.mean(discounted_epr)
109    discounted_epr /= np.std(discounted_epr)
110
111    epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112    grad = policy_backward(eph, epdlogp)
113    for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
```

注意一次游戏，有许多轮，每一轮才是一个完整的序列

4.训练 (4)

```
114
115     # perform rmsprop parameter update every batch_size episodes
116     if episode_number % batch_size == 0:
117         for k,v in model.iteritems():
118             g = grad_buffer[k] # gradient
119             rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120             model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121             grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123     # boring book-keeping
124     running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125     print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126     if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127     reward_sum = 0
128     observation = env.reset() # reset env
129     prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + ('' if reward == -1 else ' !!!!!!!')
```

- 实际Demo（视频演示）